

Supplement for RDD results

2026-09-15 08:35

Abstract

Supplemental material to support the RDD result in Arjmandi-Lari, Z., Mantzarlis, A., Stafford, T. (2025). [Threats to the sustainability of Community Notes on X](#). arXiv preprint arXiv:2510.00650

Table of contents

1	Introduction	1
1.1	Topline results and v1-v2 chanegs	2
2	Setup & Data Import	2
3	Diagnostic Plots	3
4	Trad OLS RDD	5
5	RDD using fuzzy rdrobust with a donut	6
6	Converting rate to baseline outcome rate	10
7	Showing an illustrative RDD plot	11
8	Robustness check - density test	12
9	Robustness check - Donut RDD	13
10	Sensitivity check - RDD sensitivity with different cutoffs	14
11	Sensitivity check - data for different time windows	17

1 Introduction

Find all of this in the repo https://github.com/zahra-arjm/community_notes

This PDF made by a quarto document <supplement.qmd>. The quarto file renders the outputs for the full set of RDD results (sensitivity check etc), and also pushes some outputs to /paper-outputs as a .tex fragments so they can be pulled into overleaf directly.

To see the code which is running you need to open the .qmd file. As a PDF this shows commentary (like this text) and code outputs (in text, and images).

AI declaration: all the words are human-written, all the logic and claims are human endorsed (by TS, so they may contain errors). The code was co-written with using an AI agent, some comments from the agent are left in the code (invisible in the PDF).

1.1 Topline results and v1-v2 chanegs

- Since v1 we adopt a more sophisticated RDD method, fix bugs and are stricter on best practice with respect to causal inference. There is still a positive effect, but it is not significant (confidence bounds overlap an effect of 0).
- A zero-indexing bug meant previously we were looking at the effect of first note status on 3rd note publication. This is now fixed.
- There is (remains) a robust effect for authoring a 3rd note (`have_another_note = note_counts[(note_counts >= 3)]`).
- We discover the RDD is not sharp, due to the fact that we have the score (intercept) from the data cut off point, which is NOT the intercept at the point when the note was or was not published (scores continuing to evolve post publication). This non-sharpness can be dealt with by running a fuzzy-RDD which accounts for the change in probabilities of being published at different intercept scores (there is still a discontinuity in these probailities).
- In addition there is a special population of Notes just below the cutoff threshold (intercept 0.39 - 0.40) which have a high probability of having been published. This is due to a documented *feature* of the CN system, whereby notes are protected from being unpublished when their score drops by a small amount (<0.1). These can be handled by using a “Donut” RDD, where notes +/- 0.01 of the cutoff are excluded
- To do all of this we move from using OLS to the robust RDD package
- We also run additional sensitivity checks and robustness checks (e.g. local density check around threshold)

2 Setup & Data Import

We fix the `cumsum=>2` issue, so we have authors who write 2 (not 3) notes as our outcome variable
We also remove the alignment filter, since this is conditioning on a collider (and makes our RDD artificially sharp)


```
Imported data for 3173276 notes
ALERT! re_auth_thresh = 2
Now we have 330234 first time note authors (all time)
Now we have 260997 first time note authors 2024-01-01 to 2026-06-02
Now we have 234819 notes written by our 4 choosen algorithms
Now we have 175194 notes classified as MISLEADING
Now we have 157044 notes which have >0 ratings
----- No filter for alignment
FINAL total 157044 Notes
of which author writes again = 107099
```

NB current (final), not Max, score

We are now also using Notes evaluated by `{python} f"{len(first_notes['decidedBy'].unique())}"` models

3 Diagnostic Plots


Most notes have a small number of rating, a small number have very many ratings. This means the distribution of ratings is very skewed. We take the log of the number of ratings to make the distribution more normal.



The basis for our RDD is the discontinuity between the intercept and note status (which defines publication / the possibility for publication)

Plotting $\log(\text{numRatings})$ vs NoteIntercept, with colours according to Note final status (finalRatingStatus)

You can see the fuzzy discontinuity - some $>\text{cutoff}$ notes are “NEEDS_MORE_RATINGS”



Sadly, though, our data record intercept calculated from the full set of ratings, this is not the contemporaneous intercept score which determines moment-by-moment publication. It approximates it, but with contamination. Notes may be published then unpublished as they receive more ratings and their rating degrades, or there may be additional blocks on publication (e.g. requirements on total number or type of ratings, or requirements on intercept stability or a lag before publication). The CN system is not entirely clear about all these details.

These contamination can be seen by plotting indicators of ever having been published against our (final) intercept.

We construct an independent ‘ever_published’ variable from the `firstNonNMRStatus` and `mostRecentNonNMRStatus` variables. The logic being that if either of these are `CURRENTLY_RATED_HELPFUL` then the Note must have achieved CRH/published status at some point in the past (ie separate from the `finalRatingStatus` score)


```
[0.35,0.40): n= 4979 P(ever_pub)=0.3866 P(ever_crit)=0.4832
[0.38,0.40): n= 1942 P(ever_pub)=0.5695 P(ever_crit)=0.6560
[0.39,0.40): n= 1205 P(ever_pub)=0.7137 P(ever_crit)=0.7826
[0.40,0.42): n= 2102 P(ever_pub)=0.8782 P(ever_crit)=0.9172
```

Below-cutoff `P(if_written_again)` by `ever_published`:

	mean	size
<code>ever_published</code>		
0	0.669529	104094
1	0.743547	4765

Below-cutoff `P(if_written_again)` by `ever_met_crh_criteria`:

	mean	size
<code>ever_met_crh_criteria</code>		
False	0.670453	102747
True	0.711715	6112



This shows that our secondary indicator of publication status shows a discontinuity near the cutoff, but not a sharp one

The high proportion of ‘ever published’ just below the cutoff suggests many notes got published and then slipped back below the threshold

We deal with this by a) using a fuzzy RDD and b) excluding the just below cutoff points by using the donut method

4 Trad OLS RDD

Our first approach was to use OLS and an arbitrary window

This is now significant with a wide window (0.08, it was non-significant at 0.05), and a 0.01 donut (but not significant if donut = 0, because of the contamination in the 0.39-0.40 zone))

bandwidth = 0.08

Number of treated units: 7418

Number of control units: 8419

Estimated Treatment Effect at the Cutoff: 0.0539

I think we can intepret this as a % change the chance of authoring again if published

P-value for treatment: 0.0055

T-statistic for treatment: 2.7795

This OLS approach provides a bridge to the fuzzy-RDD approach, which uses robust RDD to estimate the window size (as well as bringing in a bunch of statistical machinery which is best practice for RDDs)

5 RDD using fuzzy rdrobust with a donut

The rdrobust function will estimate the bandwidth from the data, as well as provide various sensitivity checks and allow a fuzzy RDD

It has the option of using a polynomial fit (which we decline), and uses all points, weighted by proximity to cutoff

We are using a ‘donut’ method of excluding those points near the cutoff, this is specifically to get those points just below the cutoff which have a high (outlier) percentage change of ever having been published

References for this method

@article{calonico2014robust, title = {Robust Nonparametric Confidence Intervals for Regression-Discontinuity Designs}, author = {Calonico, Sebastian and Cattaneo, Matias D. and Titiunik, Rocio}, journal = {Econometrica}, volume = {82}, number = {6}, pages = {2295–2326}, year = {2014}, doi = {10.3982/ECTA11757} }

@article{calonico2018effect, title = {On the Effect of Bias Estimation on Coverage Accuracy in Nonparametric Inference}, author = {Calonico, Sebastian and Cattaneo, Matias D. and Farrell, Max H.}, journal = {Journal of the American Statistical Association}, volume = {113}, number = {522}, pages = {767–779}, year = {2018}, doi = {10.1080/01621459.2017.1285776} }

@article{calonico2014rdrobust, title = {Robust Data-Driven Inference in the Regression-Discontinuity Design}, author = {Calonico, Sebastian and Cattaneo, Matias D. and Titiunik, Rocio}, journal = {The Stata Journal}, volume = {14}, number = {4}, pages = {909–946}, year = {2014}, doi = {10.1177/1536867X1401400413} }

Justifying the donut RDD, we can test if intercept (our running variable) is associated with ‘ever published’

PRIMARY SPEC - donut ± 0.01

Call: rdrobust

Sharp RD estimates using local polynomial regression.

Number of Observations:	122098	
Polynomial Order Est. (p):	1	
Polynomial Order Bias (q):	2	
Kernel:	Triangular	
Bandwidth Selection:	mserd	
Var-Cov Estimator:	NN	
	Left	Right

Number of Observations	107654	14444

Number of Unique Obs.	107449	14426
Number of Effective Obs.	3630	4151
Bandwidth Estimation	0.049	0.049
Bandwidth Bias	0.119	0.119
rho (h/b)	0.409	0.409

Method	Coef.	S.E.	z-stat	P> z	95% CI
Conventional	0.486	0.031	15.549	1.626e-54	[0.424, 0.547]
Bias-Corrected	0.471	0.031	15.073	2.447e-51	[0.41, 0.532]
Robust	0.471	0.034	13.654	1.902e-42	[0.403, 0.538]

- - - - variation: donut ± 0.0

Call: rdrobust

Sharp RD estimates using local polynomial regression.

Number of Observations:	124389
Polynomial Order Est. (p):	1
Polynomial Order Bias (q):	2
Kernel:	Triangular
Bandwidth Selection:	mserd
Var-Cov Estimator:	NN

	Left	Right
Number of Observations	108859	15530
Number of Unique Obs.	108653	15510
Number of Effective Obs.	1864	1973
Bandwidth Estimation	0.019	0.019
Bandwidth Bias	0.058	0.058
rho (h/b)	0.325	0.325

Method	Coef.	S.E.	z-stat	P> z	95% CI
Conventional	0.014	0.027	0.516	6.059e-01	[-0.039, 0.066]
Bias-Corrected	0.0	0.027	0.002	9.986e-01	[-0.052, 0.053]
Robust	0.0	0.028	0.002	9.987e-01	[-0.054, 0.054]

- - - - variation: donut ± 0.005

Call: rdrobust

Sharp RD estimates using local polynomial regression.

Number of Observations:	123222
Polynomial Order Est. (p):	1
Polynomial Order Bias (q):	2
Kernel:	Triangular

Bandwidth Selection: mserd
 Var-Cov Estimator: NN

	Left	Right
Number of Observations	108234	14988
Number of Unique Obs.	108028	14969
Number of Effective Obs.	1384	1646
Bandwidth Estimation	0.021	0.021
Bandwidth Bias	0.072	0.072
rho (h/b)	0.291	0.291

Method	Coef.	S.E.	z-stat	P> z	95% CI
Conventional	-0.064	0.049	-1.292	1.964e-01	[-0.16, 0.033]
Bias-Corrected	-0.09	0.049	-1.822	6.841e-02	[-0.186, 0.007]
Robust	-0.09	0.051	-1.755	7.922e-02	[-0.19, 0.01]

- - - variation: donut ± 0.02

Call: rdrobust

Sharp RD estimates using local polynomial regression.

Number of Observations: 120345
 Polynomial Order Est. (p): 1
 Polynomial Order Bias (q): 2
 Kernel: Triangular
 Bandwidth Selection: mserd
 Var-Cov Estimator: NN

	Left	Right
Number of Observations	106917	13428
Number of Unique Obs.	106712	13412
Number of Effective Obs.	4909	4713
Bandwidth Estimation	0.063	0.063
Bandwidth Bias	0.179	0.179
rho (h/b)	0.355	0.355


Method	Coef.	S.E.	z-stat	P> z	95% CI
Conventional	0.472	0.032	14.762	2.564e-49	[0.41, 0.535]
Bias-Corrected	0.45	0.032	14.046	8.199e-45	[0.387, 0.512]
Robust	0.45	0.034	13.097	3.440e-39	[0.382, 0.517]

So this is saying that if you exclude at 0 (no exclusions) or 0.05 our running variable isn't cleanly associated with "ever published" (at 0.0 it is n.s. at 0.05 it is sig, but negative). At 0.01 the effect is positive and robust (significant), so this justifies the donut

The flip in significance shows significant contamination at the cutoff (those just below *more* likely to be ever published than those just above!). So the justification for the donut is not just that it excludes problematic points, but the RDD is only possible against a background of linear stability (which would not exist without the donut).

Note also the negligible drop in N with donut 0.01. We're not losing too much data this way.

We can visualise the donut, and this in particular provides the justification (showing that our contamination of high $p(\text{ever_published})$ Notes below the cutoff is particularly high in the 0.39-0.40)



'Mephi:\nWhy: the platform has a documented hysteresis rule of exactly 0.01. I just pulled the scoring/src/scoring/mf_base_scorer.py:178,188\nAddCRHInertia rule: a note that is already CRH I i.e. [0.39, 0.40). Class in scoring_rules.py:1521, docstring literally says "notes with interce hence the sharp break your plot shows precisely at the grey band\'s left edge.\n'

So this shows the operation fo the inertiaDelta = 0.01 rule, see scoring_rules.py line 1521. Notes which are published and drop 0.01 or less keep their status <https://github.com/twitter/communitynotes/blob/>

The OLS serves as a sanity check, but now let's run the full fuzzy RDD

```
running variable (x) length 122098
binary outcome (y) length 122098
cutoff (c) is 0.4
Call: rdrobust
```

Fuzzy RD estimates using local polynomial regression.

Number of Observations: 122098

Polynomial Order Est. (p): 1
 Polynomial Order Bias (q): 2
 Kernel: Triangular
 Bandwidth Selection: mserd
 Var-Cov Estimator: NN

	Left	Right
Number of Observations	107654	14444
Number of Unique Obs.	107449	14426
Number of Effective Obs.	5256	5489
Bandwidth Estimation	0.061	0.061
Bandwidth Bias	0.147	0.147
rho (h/b)	0.416	0.416

Method	Coef.	S.E.	z-stat	P> z	95% CI
Conventional	0.076	0.059	1.284	1.990e-01	[-0.04, 0.192]
Robust	-	-	1.164	2.445e-01	[-0.052, 0.205]

First-Stage Estimates.

Method	Coef.	S.E.	z-stat	P> z	95% CI
Conventional	0.476	0.024	19.506	9.772e-85	[0.429, 0.524]
Robust	-	-	16.982	1.116e-64	[0.405, 0.511]

MSE-optimal h = 0.0610

Bias-corrected point estimate: 0.0765 95% CI [-0.0524, 0.2054]

Bandwidth sensitivity (bias-corrected point, robust 95% CI):

h = 0.0305 (0.5x) tau = +0.1352 95% CI [-0.1293, +0.4211] N_L/N_R = 1653/2148
 h = 0.0610 (1.0x) tau = +0.0758 95% CI [-0.0524, +0.2054] N_L/N_R = 5256/5489
 h = 0.1219 (2.0x) tau = +0.0841 95% CI [+0.0235, +0.1582] N_L/N_R = 17543/11177

This is positive but non-significant

6 Converting rate to baseline outcome rate

Common bandwidth h = 0.0610

Reduced form (ITT): +0.0361 robust CI [-0.0585, +0.1963]
 First stage: +0.4765 robust CI [+0.3963, +0.6206]
 Fuzzy LATE: +0.0758 robust CI [-0.0524, +0.2054]

Tie-out check: RF / FS = +0.0758 (should LATE at common h)


The RD robust method identifies a MSE-optimal bandwidth of h = 0.061), and provides an es-

estimated treatment effect of $\downarrow \hat{\tau} = +0.076$. This reflects +7.6 percentage point increase in the probability of writing a second note if an author's first note is published (conventional local-linear estimator; robust bias-corrected 95% CI [-5.2, +20.5], $p = 0.24$). Calonico et al., 2014

The rate of republishing for Notes which fall before the cutoff of ~73.4%. The treatment effect of 7.6pp, implies a 11.5% relative lift over the counterfactual baseline, however this difference is non-significant.

7 Showing an illustrative RDD plot

First, binned means with global linear fit (this is the default plot from `rdplot`)



```
['rdplot_mean_bin', 'rdplot_mean_x', 'rdplot_mean_y', 'rdplot_min_bin', 'rdplot_max_bin', 'rdplot_data']
paper-outputs/RDDrobust_scatter.png + paper-outputs/RDDrobust_scatter_data.csv written (32 bins)
```

RD plot at $c = 0.4$ ($p=1$, triangular, $N=122098$)

```
bins used      : 29 left / 4 right
J_IMSE         : [29.0, 4.0]
J_MV           : [974.0, 1000.0]
scale          : [1, 1]
bin width      : avg [0.03054120705045503, 0.07023575305938719], median [0.030541207050455044, 0.07023575305938719]
```


coefficients:

	Left	Right
0	0.706061	0.698152
1	0.119146	0.218127

wrote paper-outputs/RDDrobust_scatter_meta.json, paper-outputs/RDDrobust_scatter_coef.csv

This shows all the data, and misleads because RD robust weights data closer to the cutoff more.

Now with the local-linear fits (over the rdrobust defined window), which more accurately represents what the model is doing



consistency check: gap at cutoff should be conventional tau
plot gap at $c = +0.0361$ vs conventional tau = $+0.0758$

8 Robustness check - density test

Next, manipulation test/density smoothness at cutoff. The density test is standard for RDD, to show that nothing weird is happening in *counts* at the cutoff.

This also, incidentally, shows the donut cut


Inspecting density with a window of $\pm 0 = 0.060969791351346406$

```
t_asy      NaN
t_jk       0.945072
p_asy      NaN
p_jk       0.344622
dtype: float64
left       0.424248
right      0.577593
diff       0.153345
dtype: float64
left       0.057317
right      0.047540
dtype: float64
full       122098
left       107654
right      14444
```

```

eff_left      4704
eff_right     4001
dtype: int64
leftN          [5256.0]
rightN         [5489.0]
leftWindow     [0.060969791351346406, 0.060969791351346406, 0...
rightWindow    [0.060969791351346406, 0.060969791351346406, 0...
pval           [0.02520822086331581]
dtype: object
['X_max', 'X_min', 'binNW', 'bino', 'binoN', 'binoNStep', 'binoP', 'binoW', 'binoWStep', 'bino']

```



The density does progressively decrease up to the cutoff, but there are but no signs of sub-threshold notes being shifted to above threshold (decline not matched by leap post-cutoff)

9 Robustness check - Donut RDD

We can run the analysis with different size donuts. This isn't really necessary, since our donut size is defined from first principles (the 0.1 hysteresis on Notes getting unpublished), but let's have a look anyway).

donut	n_eff_L	n_eff_R	h	coef_conv	se_conv	p_robust	ci_robust_lo	ci_robust_hi
±0.000	2810	3161	0.030	-0.4961	0.5596	0.2020	-1.8987	0.4014
±0.005	3217	3773	0.040	0.0150	0.1745	0.9989	-0.3617	0.3612
±0.010	5256	5489	0.061	0.0758	0.0590	0.2445	-0.0524	0.2054
±0.020	5681	5220	0.068	-0.0024	0.0741	0.9352	-0.1715	0.1578

This shows that the effect never gets significant at any donut width (ie the non-effect is robust). You can see also that the donut = 0 row is aberrant (coefficient has negative sign), which is due to the contamination which justified a non-zero donut in the first place

10 Sensitivity check - RDD sensitivity with different cutoffs

Only a cutoff of 0.4 makes sense in terms of the CN publication mechanism. Therefore we expect other thresholds to show no effect (definitely below 0.4, cutoffs above 0.4 will include some published notes above and below that cutoffs, so if there is a true effect may report an attenuated form of that. We can adjust for this by running the RDD on the data above 0.4 only for tresholds above 0.4)

```
running variable (x) length 120345
binary outcome (y) length 120345
cutoff (c) is 0.3
Call: rdrobust
```

Sharp RD estimates using local polynomial regression.

```
Number of Observations:      120345
Polynomial Order Est. (p):      1
Polynomial Order Bias (q):      2
Kernel:                      Triangular
Bandwidth Selection:           mserd
Var-Cov Estimator:             NN
```

	Left	Right
Number of Observations	95319	25026
Number of Unique Obs.	95139	24985
Number of Effective Obs.	14966	9491
Bandwidth Estimation	0.058	0.058
Bandwidth Bias	0.111	0.111
rho (h/b)	0.523	0.523

Method	Coef.	S.E.	z-stat	P> z	95% CI
Conventional	0.009	0.013	0.682	4.954e-01	[-0.017, 0.034]
Robust	-	-	0.763	4.452e-01	[-0.018, 0.041]

MSE-optimal h = 0.0579

Bias-corrected point estimate: 0.0116 95% CI [-0.0182, 0.0414]

Bandwidth sensitivity (bias-corrected point, robust 95% CI):

```
h = 0.0290 (0.5x) tau = +0.0244 95% CI [-0.0119, +0.0680] N_L/N_R = 7003/5479
h = 0.0579 (1.0x) tau = +0.0089 95% CI [-0.0182, +0.0414] N_L/N_R = 14966/9491
h = 0.1159 (2.0x) tau = +0.0041 95% CI [-0.0136, +0.0296] N_L/N_R = 33791/11598
```

```
running variable (x) length 120345
binary outcome (y) length 120345
cutoff (c) is 0.35
Call: rdrobust
```

Sharp RD estimates using local polynomial regression.

```

Number of Observations:      120345
Polynomial Order Est. (p):    1
Polynomial Order Bias (q):    2
Kernel:                      Triangular
Bandwidth Selection:          mserd
Var-Cov Estimator:           NN

```

	Left	Right
Number of Observations	103880	16465
Number of Unique Obs.	103680	16444
Number of Effective Obs.	28311	8803
Bandwidth Estimation	0.124	0.124
Bandwidth Bias	0.183	0.183
rho (h/b)	0.676	0.676

Method	Coef.	S.E.	z-stat	P> z	95% CI
Conventional	-0.001	0.012	-0.099	9.209e-01	[-0.025, 0.023]
Robust	-	-	-0.288	7.732e-01	[-0.032, 0.024]

MSE-optimal h = 0.1236

Bias-corrected point estimate: -0.0041 95% CI [-0.0322, 0.0239]

Bandwidth sensitivity (bias-corrected point, robust 95% CI):

```

h = 0.0618 (0.5x) tau = +0.0053 95% CI [-0.0375, +0.0505] N_L/N_R = 11229/3037
h = 0.1236 (1.0x) tau = -0.0012 95% CI [-0.0322, +0.0239] N_L/N_R = 28311/8803
h = 0.2473 (2.0x) tau = +0.0038 95% CI [-0.0188, +0.0272] N_L/N_R = 68432/16265

```

running variable (x) length 13428

binary outcome (y) length 13428

cutoff (c) is 0.45

Call: rdrobust

Sharp RD estimates using local polynomial regression.

```

Number of Observations:      13428
Polynomial Order Est. (p):    1
Polynomial Order Bias (q):    2
Kernel:                      Triangular
Bandwidth Selection:          mserd
Var-Cov Estimator:           NN

```

	Left	Right
--	------	-------

Number of Observations	3273	10155
Number of Unique Obs.	3270	10142
Number of Effective Obs.	871	876
Bandwidth Estimation	0.008	0.008
Bandwidth Bias	0.012	0.012
rho (h/b)	0.637	0.637

Method	Coef.	S.E.	z-stat	P> z	95% CI
Conventional	0.091	0.048	1.907	5.647e-02	[-0.003, 0.185]
Robust	-	-	1.767	7.720e-02	[-0.011, 0.213]

MSE-optimal h = 0.0078

Bias-corrected point estimate: 0.1008 95% CI [-0.0110, 0.2126]

Bandwidth sensitivity (bias-corrected point, robust 95% CI):

h = 0.0039 (0.5x)	tau = +0.0611	95% CI [-0.1093, +0.2078]	N_L/N_R = 461/443
h = 0.0078 (1.0x)	tau = +0.0911	95% CI [-0.0110, +0.2126]	N_L/N_R = 871/876
h = 0.0157 (2.0x)	tau = +0.0537	95% CI [-0.0178, +0.1385]	N_L/N_R = 1723/1687

Call: rdrobust

Sharp RD estimates using local polynomial regression.

Number of Observations:	13428
Polynomial Order Est. (p):	1
Polynomial Order Bias (q):	2
Kernel:	Triangular
Bandwidth Selection:	mserd
Var-Cov Estimator:	NN

	Left	Right
Number of Observations	3273	10155
Number of Unique Obs.	3270	10142
Number of Effective Obs.	871	876
Bandwidth Estimation	0.008	0.008
Bandwidth Bias	0.012	0.012
rho (h/b)	0.637	0.637

Method	Coef.	S.E.	z-stat	P> z	95% CI
Conventional	0.091	0.048	1.907	5.647e-02	[-0.003, 0.185]
Robust	-	-	1.767	7.720e-02	[-0.011, 0.213]

All these are non-significant and the effect estimates are close to zero

11 Sensitivity check - data for different time windows

We skip this, since there is no positive result to interrogate